

DevOps and MLOps Lab Manual

Engineering Department

January 2026

Lab Questions

1. DevOps and MLOps Lab Manual

Create a simple User Registration Form for an event with fields (Name, Email, Phone). Initialize a Git repository, commit the project files, and push them to GitHub using appropriate Git commands. Create a new branch named `update-form`, modify the registration form by adding a new field (Department), merge the branch into `main`, and push the changes to GitHub. Show the use of: `git branch`, `git checkout`, `git merge`, `git push`.

2. Your college department is developing a small student portal website hosted on GitHub. As a DevOps engineer, you must set up Continuous Integration using Jenkins:

Task Requirements

Configure a Jenkins Freestyle Job that automatically pulls the website code from GitHub.

Enable Jenkins to build the project every time a change is pushed.

Make a small update to the HTML page, push it to GitHub, and show that Jenkins automatically triggers a new build.

3. Create a simple web application (HTML or Python Flask). Write a Dockerfile, build a Docker image, and run the container so the application is accessible on a browser using port mapping. Using Docker commands, perform the following operations:

List images and containers

Stop a running container

Remove a container and image

Demonstrate commands like: `docker ps`, `docker stop`, `docker rm`, `docker rmi`.

4. Deploy the previously created Docker image on Kubernetes using a Deployment YAML file. Verify that the pod is running using `kubectl` commands. Expose the application using a NodePort Service, access it using the node's IP and port, and scale the deployment to 3 replicas using `kubectl scale`.

5. Set up a simple ML project environment by creating a `requirements.txt` file, installing packages, and verifying environment setup. Document the steps in a Jupyter Notebook and commit it to the Git repository.

6. To design and deploy a multi-container application using Docker Compose by creating an application service and a dependent database/Redis service, configuring service dependencies, networks, and volumes, and verifying inter-container communication.

7. To implement data ingestion, cleaning, and versioning using Data Version Control (DVC) by tracking raw and processed datasets, creating a reproducible pipeline with `dvc.yaml`, and validating reproducibility via `dvc repro`.

8. To perform experiment tracking using MLflow by training a machine learning model, logging metrics, parameters, and artifacts, and comparing multiple runs to identify the best-performing model.

9. To optimize and standardize model inference using ONNX by exporting a trained ML model to ONNX format, running inference with ONNX Runtime, and benchmarking performance against the native scikit-learn model.

10. To serve a machine learning model using FastAPI by developing a `/predict` REST endpoint with input validation, writing test cases, containerizing the service, and verifying predictions through Postman or `curl`.

Prerequisites

Ensure the following tools are installed and properly configured on your Linux system (Ubuntu-based). Commands below assume `sudo` privileges.

Git Installation and Verification

Download & Install:

```
1 $ sudo apt update
2 Hit:1 http://archive.ubuntu.com/ubuntu noble InRelease
3 Get:2 http://archive.ubuntu.com/ubuntu noble-updates InRelease [128 kB]
4 ...
5 Reading package lists... Done
6 $ sudo apt install -y git
7 Reading package lists... Done
8 Building dependency tree... Done
9 The following NEW packages will be installed:
10   git git-man ...
11 After this operation, 45.3 MB of additional disk space will be used.
12 Do you want to continue? [Y/n] y
13 ...
14 Setting up git (1:2.43.0-1ubuntu...) ...
```

Verify Version:

```
1 $ git --version
2 git version 2.43.0
```

Docker & Docker Compose

Download & Install Docker Engine:

```
1 $ sudo apt update
2 $ sudo apt install -y docker.io
3 Reading package lists... Done
4 ...
5 Setting up docker.io (24.0.0-0ubuntu...) ...
6 Adding group 'docker' (GID 998) ...
```

Enable & Start Docker Service:

```
1 $ sudo systemctl enable docker
2 Synchronizing state of docker.service with SysV service script...
3 $ sudo systemctl start docker
4 $ sudo systemctl status docker
5     docker.service - Docker Application Container Engine
6       Loaded: loaded (/lib/systemd/system/docker.service; enabled)
7         Active: active (running) since Tue 2026-01-06 10:15:00 IST; 5s ago
8 ...
```

Install Docker Compose Plugin (if needed):

```
1 $ sudo apt install -y docker-compose
2 Reading package lists... Done
3 ...
4 Setting up docker-compose (1.29.x-...) ...
```

Verify Docker:

```
1 $ docker --version
2 Docker version 24.0.x, build abcdefg
3 $ docker-compose --version
4 docker-compose version 1.29.x, build hijklmno
```

Jenkins

Download & Install Jenkins (Debian package method summarized):

```
1 $ curl -fsSL https://pkg.jenkins.io/debian-stable/jenkins.io.key \
2   | sudo tee /usr/share/keyrings/jenkins-keyring.asc > /dev/null
3 $ echo deb [signed-by=/usr/share/keyrings/jenkins-keyring.asc] \
4   https://pkg.jenkins.io/debian-stable binary/ | \
5   sudo tee /etc/apt/sources.list.d/jenkins.list > /dev/null
6 $ sudo apt update
7 $ sudo apt install -y openjdk-17-jre-headless jenkins
8 Reading package lists... Done
9 ...
10 Setting up jenkins (2.xxx) ...
```

Start Jenkins and Verify:

```
1 $ sudo systemctl enable jenkins
2 $ sudo systemctl start jenkins
3 $ sudo systemctl status jenkins
4     jenkins.service - Jenkins Continuous Integration Server
5       Loaded: loaded (/lib/systemd/system/jenkins.service; enabled)
6         Active: active (running) since Tue 2026-01-06 10:20:00 IST; 10s ago
7 ...
...
```

Access Jenkins in a browser at:

```
1 http://localhost:8080
```

Kubernetes (Minikube Example)

Download Minikube Binary:

```
1 $ curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
2   % Total    % Received % Xferd  Average Speed   Time   Time     Current
3                               Dload  Upload   Total Spent  Left  Speed
4 100  74.4M  100  74.4M    0      0  8.5M    0  0:00:08  0:00:08  --:--:--  9.0M
5 $ sudo install minikube-linux-amd64 /usr/local/bin/minikube
```

Start a Local Cluster:

```
1 $ minikube start
2 * minikube v1.34.0 on Ubuntu 24.04
3 * Using the docker driver based on existing profile
4 * Starting control plane node minikube in cluster minikube
5 ...
6 * Done! kubectl is now configured to use "minikube" cluster and "default" namespace
```

Verify kubectl (Minikube bundles it or install via apt):

```
1 $ kubectl get nodes
2 NAME      STATUS   ROLES      AGE      VERSION
3 minikube  Ready    control-plane  5m1s   v1.29.x
```

Python, pip, venv, and Jupyter

Install Python 3 and pip:

```
1 $ sudo apt update
2 $ sudo apt install -y python3 python3-pip python3-venv
3 Reading package lists... Done
4 ...
5 Setting up python3 (3.12.x....) ...
```

Create Virtual Environment:

```

1 $ python3 -m venv ml-env
2 $ source ml-env/bin/activate
3 (ml-env) $ python --version
4 Python 3.12.x

```

Install Jupyter and ML Packages:

```

1 (ml-env) $ python -m pip install --upgrade pip
2 (ml-env) $ pip install jupyter pandas scikit-learn mlflow fastapi uvicorn dvc
   onnxruntime
3 Collecting jupyter
4 Collecting pandas
5 Collecting scikit-learn
6 ...
7 Successfully installed jupyter-1.x pandas-2.x scikit-learn-1.x mlflow-2.x fastapi-0.11x
   ...

```

Launch Jupyter Notebook:

```

1 (ml-env) $ jupyter notebook
2 [I 10:30:00.000 NotebookApp] Serving notebooks from local directory: /home/user/project
3 [I 10:30:00.001 NotebookApp] The Jupyter Notebook is running at:
4 [I 10:30:00.001 NotebookApp] http://localhost:8888/?token=...

```

1 Git Initialization & GitHub Push

Objective: Create a form and push it to a remote repository.

File: index.html

```

1 <!DOCTYPE html>
2 <html>
3 <head><title>Event Registration</title></head>
4 <body>
5   <h2>Registration Form</h2>
6   <form>
7     Name: <input type="text"><br>
8     Email: <input type="email"><br>
9     Phone: <input type="text"><br>
10    <button>Submit</button>
11  </form>
12 </body>
13 </html>

```

Terminal Commands & Output

```

1 $ git init
2 Initialized empty Git repository in /home/user/lab1/.git/
3 $ git add index.html
4 $ git commit -m "Initial commit"
5 [main (root-commit) 4e2a1b] Initial commit
6   1 file changed, 11 insertions(+)
7   create mode 100644 index.html
8 $ git remote add origin https://github.com/username/event-form.git
9 $ git push -u origin main
10 Enumerating objects: 3, done.
11 Counting objects: 100% (3/3), done.
12 Writing objects: 100% (3/3), 245 bytes | 245.00 KiB/s, done.
13 To https://github.com/username/event-form.git
   * [new branch]      main -> main
14 branch 'main' set up to track 'origin/main'.
15

```

2 Git Branching, Merging & Push

Objective: Add a “Department” field using a feature branch.

Terminal Commands & Output

```
1 $ git branch update-form
2 $ git checkout update-form
3 Switched to branch 'update-form'
4 # (Modify index.html to add <input type="text" name="Dept">)
5 $ git add index.html
6 $ git commit -m "Added department field"
7 [update-form 7a2b3c] Added department field
8 1 file changed, 1 insertion(+)
9 $ git checkout main
10 Switched to branch 'main'
11 $ git merge update-form
12 Updating 4e2a1b..7a2b3c
13 Fast-forward
14   index.html | 1 +
15 1 file changed, 1 insertion(+)
16 $ git push origin main
17 Enumerating objects: 5, done.
18 Counting objects: 100% (5/5), done.
19 Writing objects: 100% (3/3), 280 bytes | 280.00 KiB/s, done.
20 To https://github.com/username/event-form.git
21     4e2a1b..7a2b3c  main -> main
```

3 Jenkins CI Integration

Objective: Automatically trigger builds on push.

Terminal Commands & Output

```
1 $ echo " " >> index.html
2 $ git status
3 On branch main
4 Changes not staged for commit:
5   (use "git add <file>..." to update what will be committed)
6   modified:   index.html
7
8 $ git commit -am "Trigger Jenkins"
9 [main 9c1d2e3] Trigger Jenkins
10 1 file changed, 1 insertion(+)
11 $ git push origin main
12 To https://github.com/username/event-form.git
13     7a2b3c..9c1d2e3  main -> main
```

Jenkins Console Log

```
1 Started by GitHub push by username
2 Running as SYSTEM
3 Building in workspace /var/lib/jenkins/workspace/StudentPortal
4 > git rev-parse --is-inside-work-tree # timeout=10
5 Fetching changes from the remote Git repository
6 > git fetch --tags --progress -- origin +refs/heads/*:refs/remotes/origin/*
7 > git checkout -f origin/main
8 + echo "Starting Build..."
9 Starting Build...
10 + ls
11 index.html Jenkinsfile ...
12 Finished: SUCCESS
```

4 Dockerization of Web App

Objective: Containerize a Python Flask application.

File: app.py

```
1 from flask import Flask
2 app = Flask(__name__)
3
4 @app.route('/')
5 def hello():
6     return "Portal Running!"
7
8 if __name__ == "__main__":
9     app.run(host='0.0.0.0', port=5000)
```

File: Dockerfile

```
1 FROM python:3.9-slim
2 RUN pip install flask
3 COPY app.py .
4 CMD ["python", "app.py"]
```

Terminal Commands & Output

```
1 $ docker build -t my-app .
2 Sending build context to Docker daemon 6.144kB
3 Step 1/4 : FROM python:3.9-slim
4    --> a1b2c3d4e5f6
5 Step 2/4 : RUN pip install flask
6    --> Running in 123456789abc
7 Collecting flask
8 ...
9 Successfully built 0f1e2d3c4b5a
10 Successfully tagged my-app:latest
11
12 $ docker run -d -p 8080:5000 --name web-container my-app
13 1a2b3c4d5e6f7g8h9i0j
14
15 $ docker ps
16 CONTAINER ID IMAGE COMMAND STATUS PORTS
17 1a2b3c4d5e6f my-app "python app.py" Up 10 seconds 0.0.0.0:8080->5000/tcp
18
19 $ curl http://localhost:8080
20 Portal Running!
```

5 Docker Operations

Objective: Demonstrate image and container management.

Terminal Commands & Output

```
1 $ docker ps
2 CONTAINER ID IMAGE STATUS PORTS NAMES
3 1a2b3c4d5e6f my-app Up 2 minutes 0.0.0.0:8080->5000/tcp web-container
4
5 $ docker stop web-container
6 web-container
7
8 $ docker rm web-container
```

```

9 web-container
10
11 $ docker images
12 REPOSITORY      TAG      IMAGE ID      CREATED      SIZE
13 my-app          latest   0f1e2d3c4b5a   1 minute ago   125MB
14
15 $ docker rmi my-app
16 Untagged: my-app:latest
17 Deleted: sha256:0f1e2d3c4b5a...

```

6 Kubernetes Deployment & Scaling

Objective: Deploy to Kubernetes and scale to 3 replicas.

File: deployment.yaml

```

1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: ml-deploy
5 spec:
6   replicas: 1
7   selector:
8     matchLabels:
9       app: ml
10  template:
11    metadata:
12      labels:
13        app: ml
14    spec:
15      containers:
16        - name: ml-container
17          image: my-app
18          ports:
19            - containerPort: 5000

```

Terminal Commands & Output

```

1 $ kubectl apply -f deployment.yaml
2 deployment.apps/ml-deploy created
3
4 $ kubectl get pods
5 NAME                  READY   STATUS    RESTARTS   AGE
6 ml-deploy-7f8d9-abc1   1/1     Running   0          20s
7
8 $ kubectl expose deployment ml-deploy --type=NodePort --port=5000
9 service/ml-deploy exposed
10
11 $ kubectl get svc ml-deploy
12 NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
13 ml-deploy   NodePort   10.96.123.45   <none>           5000:31456/TCP   30s
14
15 # Access via http://<node-ip>:31456
16
17 $ kubectl scale deployment ml-deploy --replicas=3
18 deployment.apps/ml-deploy scaled
19
20 $ kubectl get pods
21 NAME                  READY   STATUS    RESTARTS   AGE
22 ml-deploy-7f8d9-abc1   1/1     Running   0          1m
23 ml-deploy-7f8d9-abc2   1/1     Running   0          10s
24 ml-deploy-7f8d9-abc3   1/1     Running   0          10s

```

7 ML Environment Setup

Objective: Set up environment and document via Jupyter.

File: requirements.txt

```
1 pandas
2 scikit-learn
3 mlflow
4 jupyter
```

Terminal Commands & Output

```
1 (ml-env) $ pip install -r requirements.txt
2 Collecting pandas
3 Collecting scikit-learn
4 Collecting mlflow
5 Collecting jupyter
6 ...
7 Successfully installed jupyter-1.x mlflow-2.x pandas-2.x scikit-learn-1.x
8
9 (ml-env) $ jupyter notebook
10 [I 10:30:00.000 NotebookApp] Serving notebooks from local directory: /home/user/project
11 [I 10:30:00.001 NotebookApp] The Jupyter Notebook is running at:
12 [I 10:30:00.001 NotebookApp] http://localhost:8888/?token=...
```

8 Docker Compose (Multi-Container)

Objective: Launch web and Redis services.

File: docker-compose.yml

```
1 services:
2   app:
3     build: .
4     ports: ["5000:5000"]
5   redis:
6     image: "redis:alpine"
```

Terminal Commands & Output

```
1 $ docker-compose up -d
2 Creating network "lab_default" with the default driver
3 Creating lab_redis_1 ... done
4 Creating lab_app_1 ... done
5
6 $ docker-compose ps
7      NAME            COMMAND           STATUS        PORTS
8      lab_app_1      "python app.py"    running       0.0.0.0:5000->5000/tcp
9      lab_redis_1    "docker-entrypoint" running       6379/tcp
```

9 DVC & MLflow Tracking

Objective: Track data changes and log experiments.

Terminal Commands & Output

```
1 (ml-env) $ dvc init
2 Initialized DVC repository.
3
4 (ml-env) $ dvc add raw_data.csv
5 Adding...
6 100% Add|  

7 | 1/1 [00:00, 2.00file/s]
8 To track the changes with git, run:  

9  

10 git add raw_data.csv.dvc .gitignore  

11  

12 (ml-env) $ git add raw_data.csv.dvc .gitignore
13 (ml-env) $ git commit -m "Track data v1"
14 [main 1a2b3c4] Track data v1
15 2 files changed, 5 insertions(+)
16 create mode 100644 raw_data.csv.dvc  

17  

18 (ml-env) $ dvc repro
19 Data is up to date.  

20  

21 (ml-env) $ python train.py
22 Model trained. Accuracy: 0.92
23 Logging to MLflow...
24  

25 (ml-env) $ mlflow ui
26 INFO [mlflow.server] Starting gunicorn 20.x
INFO [mlflow.server] Listening at: http://127.0.0.1:5000 (pid=1234)
```

10 FastAPI Serving & ONNX Optimization

Objective: REST API for inference and model optimization.

Terminal Commands & Output

```
1 (ml-env) $ uvicorn main:app --reload
2 INFO:     Will watch for changes in these directories: ['/home/user/project']
3 INFO:     Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
4
5 $ curl -X POST "http://127.0.0.1:8000/predict" \
6   -H "Content-Type: application/json" \
7   -d '{"feature1":5, "feature2":10}'
8 {"prediction": 15}
9
10 (ml-env) $ python convert_onnx.py
11 Converting model to ONNX...
12 Model converted successfully. Saved as 'model.onnx'.
13
14 (ml-env) $ ls -lh model.onnx
15 -rw-r--r-- 1 user user 45K Jan 6 11:00 model.onnx
```